

# HW1


February 20, 2022

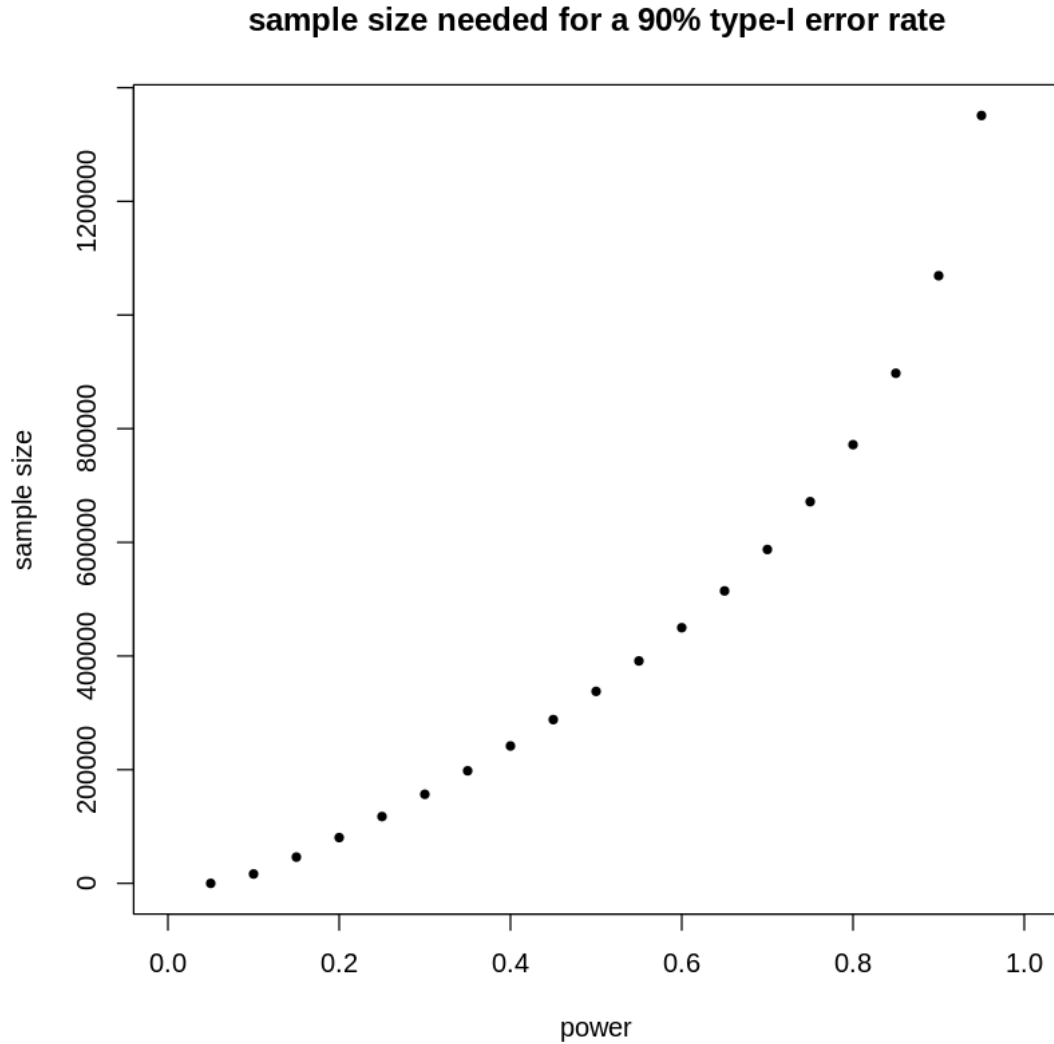
## 1 Problem 1: A/B Testing

Since currently, the company generates an average of \$1.50 per visit, so  $\mu_0 = 1.5$ . The company decideds that change in revenue with a 1% change in either direction is significant enough to note, so  $\mu_1 = 1.5 (1.01) = 1.515$ . The standard deviation  $\sigma = 5.3$ . The effect size is calculated by  $ES = \frac{|\mu_1 - \mu_0|}{\sigma} = \frac{|1.515 - 1.5|}{5.3} = 0.0028$ . Since we are assuming a 95% type-I error rate, the sample size required is  $n = \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES}\right)^2 = \frac{1.96 + Z_{1-\beta}}{0.0028}$ . Here,  $1 - \beta$  is the power of our test.

```
[2]: ES = (1.515 - 1.5) / 5.3
power = seq(0, 1, 0.05)
sample_size_95 = ((qnorm(1 - 0.05/2) + qnorm(power))/ES)^2
plot(power, sample_size_95, ylab = "sample size", main = "sample size needed_
  ↳for a 95% type-I error rate", pch = 20)
sample_size_90 = ((qnorm(1 - 0.1/2) + qnorm(power))/ES)^2
plot(power, sample_size_90, ylab = "sample size", main = "sample size needed_
  ↳for a 90% type-I error rate", pch = 20)
```

sample size needed for a 95% type-I error rate





From two scatter plots, we can see that sample size needed for a 95% type-I error rate is much higher than a 90% type-I error rate with respect to the same power of the test. For 95% type-I error rate, I will recommend a sample size (visits) around 1300000 to ensure 90% power. For 90% type-I error rate, I will recommend a sample size (visits) around 1000000 to ensure 90% power.

Power is defined as the probability of rejecting null hypothesis given the alternative hypothesis is true, which is represented as  $1 - \beta$ . Type-I error rate is  $\alpha$ . As  $\alpha$  increase,  $Z_{1-\alpha/2}$  will decrease, so the sample size will decrease. As  $\beta$  increase,  $Z_{1-\beta}$  will decrease, so the sample size will decrease as well. As effect size increase, the sample size required will decrease. As variance increase, the effect size will decrease, so the sample size required will increase.

## 2 Problem 2: Missing Data Patterns

```
[3]: data1 = read.csv("color_test.csv")
      head(data1)
      attach(data1)
```

A data.frame: 6 × 6

	id	treat	est_age	past_cust	past_avg_rev	obs_rev
	<int>	<int>	<dbl>	<int>	<dbl>	<dbl>
1	1	0	36.35849	0	NA	0.00000
2	2	0	31.47492	1	1.553581	0.00000
3	3	1	40.50781	0	NA	0.00000
4	4	0	36.68641	1	1.915854	0.00000
5	5	0	NA	1	0.000000	11.94924
6	6	0	40.05534	1	2.377523	0.00000

```
[4]: install.packages('naniar')
      library(naniar)
```


Installing package into ‘/usr/local/lib/R/site-library’  
(as ‘lib’ is unspecified)

```
[5]: sum(!complete.cases(data1)) # 7843 rows have missing data
      nrow(data1) # total number of rows is 15000
```

7843

15000

```
[6]: vis_miss(data1)
```



```
[20]: data1 = na.omit(data1) # after omitting the rows containing missing values,
      ↪ there are 7157 rows left
      nrow(data1)
      head(data1)
```

7157

		id	treat	est_age	past_cust	past_avg_rev	obs_rev
		<int>	<int>	<dbl>	<int>	<dbl>	<dbl>
	2	2	0	31.47492	1	1.553581	0
	4	4	0	36.68641	1	1.915854	0
	6	6	0	40.05534	1	2.377523	0
	7	7	0	40.05415	1	4.292266	0
	9	9	1	45.25272	1	0.000000	0
	10	10	1	35.90415	1	0.000000	0

A data.frame: 6 × 6

From above graph, we can see that only columns `est_age`, `past_avg_rev`, and `obs_rev` contain missing data. `Est_age` has 20.35% missing data, `past_avg_rev` has 34.06% missing data, and `obs_rev` has 8.01% missing data. `Est_age` is MAR because it is estimated based on other variables. `Past_avg_rev` is mar because it represents past average revenue. When `past_cust` is 0 (a new customer), it is definitely NA. `Obs_rev` is MCAR because there is no relationship between the missingness of the data and any values, observed or missing.

For this experiment, we can conduct complete-case analysis since after omitting all the rows containing missing values, we still have 7157 rows left, although this method may introduce bias to our analysis.

```
[22]: library(tidyverse)
```

```
Warning message in system("timedatectl", intern = TRUE):
```

```
"running command 'timedatectl' had status 1"
```

```
Attaching packages                                tidyverse
1.3.1
```

```
tibble 3.1.6      purrr  0.3.4
tidyr  1.1.4      stringr 1.4.0
readr  2.1.1      forcats 0.5.1
```

```
Conflicts
```

```
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()     masks stats::lag()
```

```
[ ]: women_weight <- genderweight %>%
  filter(group == "F") %>%
  pull(weight)
men_weight <- genderweight %>%
  filter(group == "M") %>%
  pull(weight)
# Compute t-test
res <- t.test(women_weight, men_weight)
res
```

```
[26]: treatment = data1 %>% filter(treat == 1) %>% select(obs_rev)
non_treatment = data1 %>% filter(treat == 0) %>% select(obs_rev)
```

```
t_test <- t.test(treatment, non_treatment)
t_test
```

Welch Two Sample t-test

```
data: treatment and non_treatment
t = 9.2936, df = 6765.6, p-value < 2.2e-16
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.8975514 1.3774129
sample estimates:
mean of x mean of y
 3.567782  2.430300
```

Since the p-value is super small, we can reject the null hypothesis that  $treatment = nontreatment$ . In other words, there is enough evidence to state that the mean observed revenue for treatment group is different to that of the non-treatment group. The difference for the mean control and test revenue is  $3.567782 - 2.4303 = 1.137482$ .

### 3 Problem 3: Exploratory Analysis

```
[8]: data2 = read.csv("US_Accidents_Dec20_updated.csv")
```

```
[9]: install.packages('dplyr')
      library(dplyr)
```

Installing package into ‘/usr/local/lib/R/site-library’  
(as ‘lib’ is unspecified)

Attaching package: ‘dplyr’

The following objects are masked from ‘package:stats’:

filter, lag

The following objects are masked from ‘package:base’:

intersect, setdiff, setequal, union

```
[10]: head(data2)
      summary(data2)
```

```
names(data2)
attach(data2)
```

```

      ID      Severity Start_Time      End_Time      Start_Lat Start_
      <chr>      <int>      <chr>      <chr>      <dbl>      <dbl>
1 A-2716600      3      2016-02-08 00:37:08 2016-02-08 06:37:08 40.10891 -83.092
2 A-2716601      2      2016-02-08 05:56:20 2016-02-08 11:56:20 39.86542 -84.062
3 A-2716602      2      2016-02-08 06:15:39 2016-02-08 12:15:39 39.10266 -84.524
4 A-2716603      2      2016-02-08 06:15:39 2016-02-08 12:15:39 39.10148 -84.523
5 A-2716604      2      2016-02-08 06:51:45 2016-02-08 12:51:45 41.06213 -81.537
6 A-2716605      3      2016-02-08 07:53:43 2016-02-08 13:53:43 39.17239 -84.492

```

```

      ID      Severity      Start_Time      End_Time
Length:1316687 Min.      :1.0 Length:1316687 Length:1316687
Class :character 1st Qu.:2.0 Class :character Class :character
Mode  :character Median :2.0 Mode  :character Mode  :character
Mean   :2.2
3rd Qu.:2.0
Max.   :4.0

```

```

      Start_Lat      Start_Lng      End_Lat      End_Lng
Min.      :24.57 Min.      : -124.50 Min.      :24.57 Min.      : -124.50
1st Qu.:33.85 1st Qu.: -118.24 1st Qu.:33.85 1st Qu.: -118.24
Median :37.15 Median : -94.60 Median :37.15 Median : -94.60
Mean   :36.80 Mean   : -98.89 Mean   :36.80 Mean   : -98.89
3rd Qu.:40.67 3rd Qu.: -80.88 3rd Qu.:40.67 3rd Qu.: -80.87
Max.   :49.00 Max.   : -67.11 Max.   :49.08 Max.   : -67.11

```

```

      Distance.mi.      Description      Number      Street
Min.      : 0.0000 Length:1316687 Min.      :      1 Length:1316687
1st Qu.: 0.0000 Class :character 1st Qu.: 1215 Class :character
Median : 0.1430 Mode  :character Median : 3999 Mode  :character
Mean   : 0.5711 Mean   : 8802
3rd Qu.: 0.5730 3rd Qu.: 10035
Max.   :155.1860 Max.   :961043
NA's   :890371

```

```

      Side      City      County      State
Length:1316687 Length:1316687 Length:1316687 Length:1316687
Class :character Class :character Class :character Class :character
Mode  :character Mode  :character Mode  :character Mode  :character

```

```

      Zipcode      Country      Timezone      Airport_Code
Length:1316687 Length:1316687 Length:1316687 Length:1316687
Class :character Class :character Class :character Class :character
Mode  :character Mode  :character Mode  :character Mode  :character

```



Weather_Timestamp	Temperature.F.	Wind_Chill.F.	Humidity...
Length:1316687	Min. : -89.00	Min. : -89.0	Min. : 1.00
Class :character	1st Qu.: 47.00	1st Qu.: 41.0	1st Qu.: 49.00
Mode :character	Median : 60.00	Median : 57.0	Median : 68.00
	Mean : 59.08	Mean : 54.8	Mean : 65.01
	3rd Qu.: 73.00	3rd Qu.: 70.0	3rd Qu.: 85.00
	Max. : 170.60	Max. : 113.0	Max. : 100.00
	NA's : 38428	NA's : 351167	NA's : 40748
Pressure.in.	Visibility.mi.	Wind_Direction	Wind_Speed.mph.
Min. : 0.00	Min. : 0.00	Length:1316687	Min. : 0.00
1st Qu.:29.41	1st Qu.: 10.00	Class :character	1st Qu.: 3.50
Median :29.87	Median : 10.00	Mode :character	Median : 7.00
Mean :29.55	Mean : 9.09		Mean : 7.49
3rd Qu.:30.04	3rd Qu.: 10.00		3rd Qu.: 10.40
Max. :58.04	Max. :130.00		Max. :984.00
NA's :31872	NA's :38793		NA's :103864
Precipitation.in.	Weather_Condition	Amenity	Bump
Min. : 0	Length:1316687	Length:1316687	Length:1316687
1st Qu.: 0	Class :character	Class :character	Class :character
Median : 0	Mode :character	Mode :character	Mode :character
Mean : 0			
3rd Qu.: 0			
Max. :24			
NA's :398289			
Crossing	Give_Way	Junction	No_Exit
Length:1316687	Length:1316687	Length:1316687	Length:1316687
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
Railway	Roundabout	Station	Stop
Length:1316687	Length:1316687	Length:1316687	Length:1316687
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
Traffic_Calming	Traffic_Signal	Turning_Loop	Sunrise_Sunset
Length:1316687	Length:1316687	Length:1316687	Length:1316687
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

Civil_Twilight	Nautical_Twilight	Astronomical_Twilight
Length:1316687	Length:1316687	Length:1316687
Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character

1. 'ID' 2. 'Severity' 3. 'Start\_Time' 4. 'End\_Time' 5. 'Start\_Lat' 6. 'Start\_Lng' 7. 'End\_Lat'  
 8. 'End\_Lng' 9. 'Distance.mi.' 10. 'Description' 11. 'Number' 12. 'Street' 13. 'Side' 14. 'City'  
 15. 'County' 16. 'State' 17. 'Zipcode' 18. 'Country' 19. 'Timezone' 20. 'Airport\_Code'  
 21. 'Weather\_Timestamp' 22. 'Temperature.F.' 23. 'Wind\_Chill.F.' 24. 'Humidity..' 25. 'Pres-  
 sure.in.' 26. 'Visibility.mi.' 27. 'Wind\_Direction' 28. 'Wind\_Speed.mph.' 29. 'Precipitation.in.'  
 30. 'Weather\_Condition' 31. 'Amenity' 32. 'Bump' 33. 'Crossing' 34. 'Give\_Way' 35. 'Junction'  
 36. 'No\_Exit' 37. 'Railway' 38. 'Roundabout' 39. 'Station' 40. 'Stop' 41. 'Traffic\_Calming' 42. 'Traf-  
 fic\_Signal' 43. 'Turning\_Loop' 44. 'Sunrise\_Sunset' 45. 'Civil\_Twilight' 46. 'Nautical\_Twilight'  
 47. 'Astronomical\_Twilight'


```
[11]: sum(!complete.cases(data2)) # 999872 rows have missing data
      nrow(data2)
      gg_miss_var(data2, show_pct = TRUE)
```

999872

1316687

Warning message:

"It is deprecated to specify `guide = FALSE` to remove a guide. Please use  
 `guide = "none"` instead."



From above graph, we can see that almost 70% of the data in column number is missing; 33% of data in column percipitation.in. is missing; 30% of data in column wind\_chill.f. is missing; 8% of data in column wind\_speed.mph. is missing. There are total eight columns contain missing values. The other four columns don't contain a lot of missing data, roughly around 3%.

```
[12]: number_cities = nrow(unique(select(data2, City)))
number_cities # there are total 10658 unique cities in the dataset
top20 = data2 %>% count(City, sort = TRUE) %>% slice(1:20)
print(top20)


barplot(height = top20$n, names.arg = top20$City, las = 2, cex.axis = 0.7, cex.
↪names = 0.8, main = "top 20 city that has highest number of accidents in the_
↪US")
```

```
title(ylab = "number of accidents")
```

10198

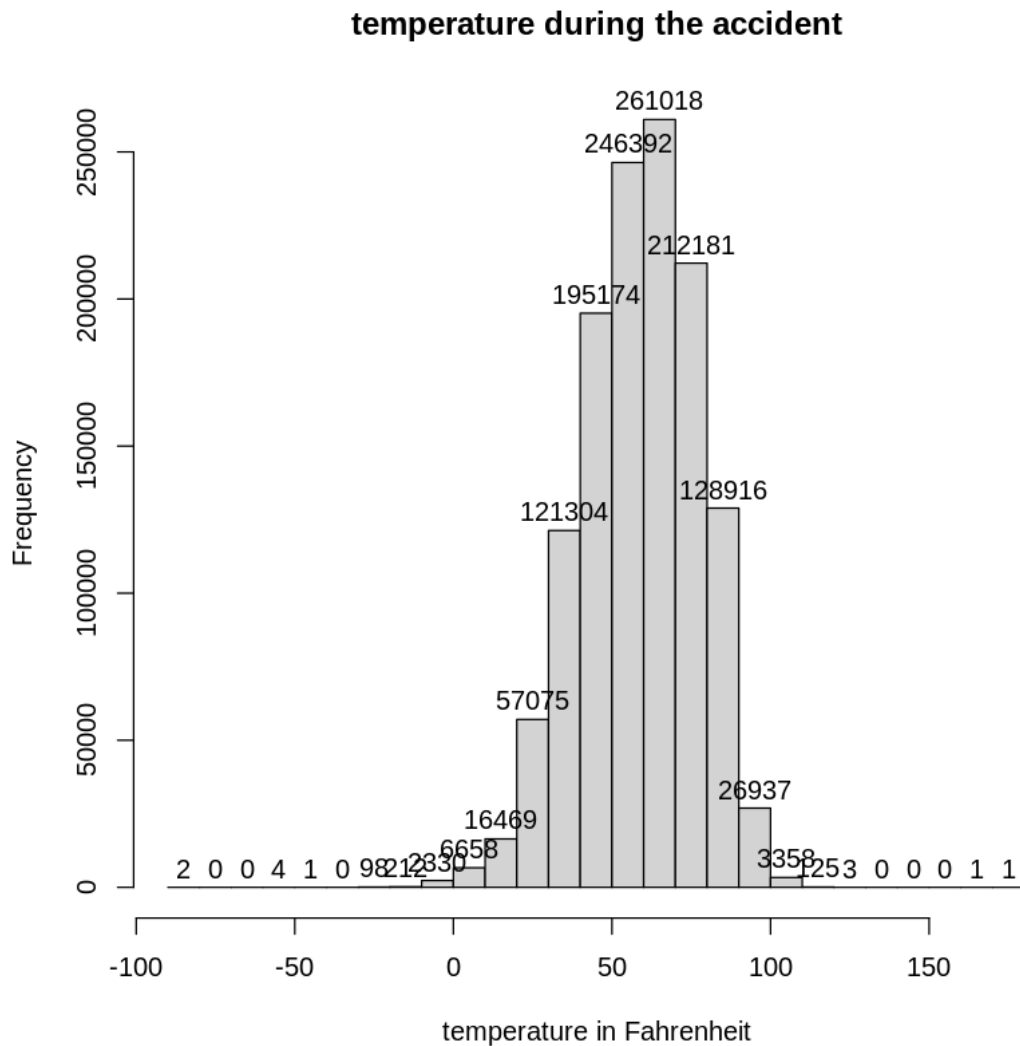
	City	n
1	Los Angeles	35818
2	Miami	33139
3	Charlotte	19692
4	Sacramento	17122
5	Houston	16827
6	Dallas	16336
7	Orlando	14324
8	San Diego	12583
9	Portland	12309
10	Minneapolis	11451
11	Saint Paul	9219
12	Atlanta	9110
13	Raleigh	8719
14	Chicago	8474
15	Phoenix	8392
16	San Jose	8117
17	Richmond	7848
18	Riverside	7481
19	Nashville	7439
20	Baton Rouge	6856

top 20 city that has highest number of accidents in the US



The above barplot shows the top 20 city in the US that has highest number of accidents. Both Los Angeles and Miami have extremely high number of accidents, more than 30000 cases. City in the third place even doesn't exceed 20000 cases. All the cities after Miami experience a steady decrease in number of accidents.

```
[13]: hist(Temperature.F., breaks = 20, main = "temperature during the accident",
→xlab = "temperature in Fahrenheit", labels = TRUE)
```



Above graph shows the local temperature when the accident happened. We can see the most common temperature range during accident is from 60 to 70 Fahrenheit.

```
[14]: #number_cities = nrow(unique(select(data2, City)))
#number_cities # there are total 10658 unique cities in the dataset
#top20 = data2 %>% count(City, sort = TRUE) %>% slice(1:20)
#print(top20)
weather = nrow(unique(select(data2, Weather_Condition)))
weather # there are total 111 unique weather situation
top10 = data2 %>% count(Weather_Condition, sort = TRUE) %>% slice(1:10)
top10
```

```

barplot(height = top10$n, names.arg = top10$Weather_Condition, las = 2, cex.axis =
  ↪= 0.7, cex.names = 0.8, main = "top 10 weather condition when there is an ↪
  ↪accident in the US")
title(ylab = "Frequency")


```

111

A data.frame: 10 × 2

Weather_Condition	n
<chr>	<int>
Fair	432670
Mostly Cloudy	164722
Cloudy	149902
Clear	136277
Partly Cloudy	111444
Overcast	66749
Light Rain	63259
	38311
Scattered Clouds	35874
Light Snow	19161

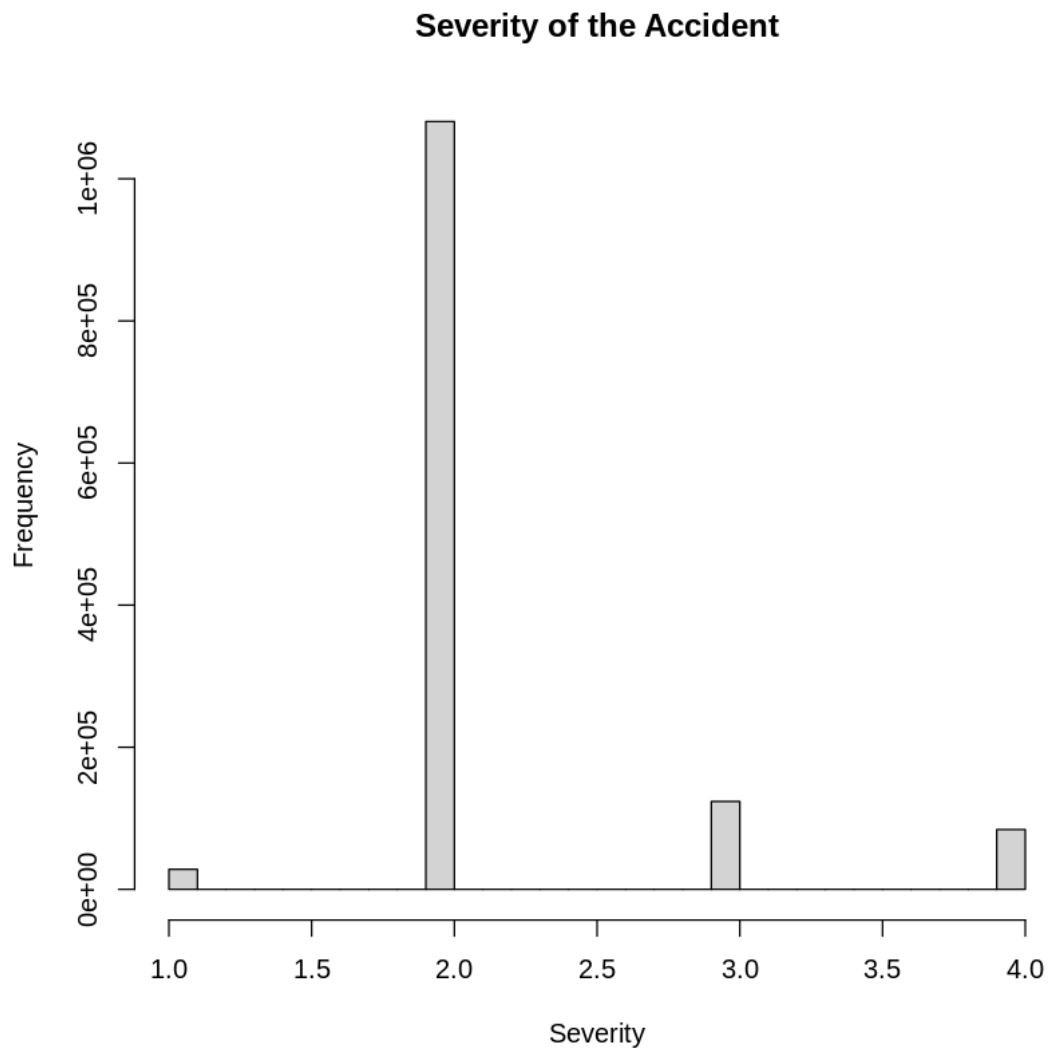
**top 10 weather condition when there is an accident in the US**



From above graph, we can see that there are more than 400000 times when the accident happened and the weather condition is fair. After that, weather condition associates with the second highest and third highest accident rate is mostly cloudy and cloudy. Fair condition contributes the most compared to others and has extremely high frequency.

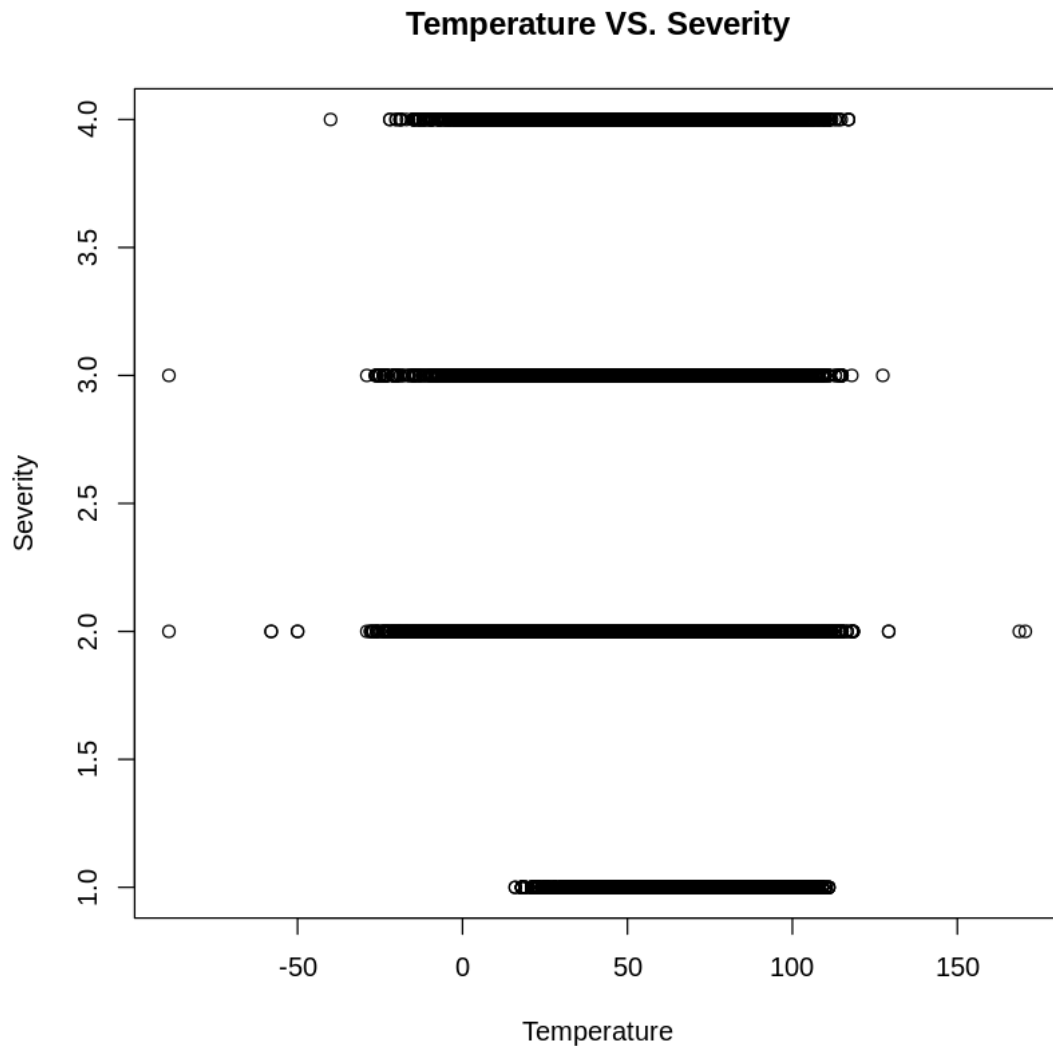
```
[15]: hist(Severity, main = "Severity of the Accident")
```





Above graph show the severity of all the accidents. One indicates the least impact and four indicates a significant impact. We can see that for most of the cases, the severity is 2, meaning that most of the accidents is somewhat severe, but not that severe.


```
[16]: plot(x=Temperature.F., y=Severity, xlab="Temperature", ylab="Severity", main = "Temperature VS. Severity")
```



From above plot, we can see that temperature ranges from -20F to 110F all have led to a car accident with severity from 2 to 4. The only exception is that when the severity is 1, the temperature associates with is from approximately 15F to 110F.

```
[17]: library(ggplot2)
```

```
[19]: p = ggplot(data = data2)
p = p + geom_point(aes(x=State, y=Distance.mi., group=Severity, color=Severity))
p = p + xlab("State") + ylab("Distance")
p
```



From above graph, we can see that in each state, the length of the road extent affected by the accident (distance) is distributed almost the same. Most of them are clustered around range 0 to 25. Moreover, the severity of the accidents also doesn't depend on which state you are in and the road extent affected by the accident. The severity is pretty much mixed up in all the states.

```
[28]: number_states = nrow(unique(select(data2, State)))
      number_states # there are total 49 unique states in the dataset
```

49

```
[34]: top20_state = data2 %>% count(State, sort = TRUE) %>% slice(1:20)
      print(top20_state)
```


```

barplot(height = top20_state$n, names.arg = top20_state$State, las = 2, cex.axis =
  ↪ 0.7, cex.names = 0.8, main = "Top 20 states that has the highest number of ↪
  ↪accidents in the US",
ylab = "Number of Accidents", xlab = "State")

```

	State	n
1	CA	408771
2	FL	136856
3	OR	74653
4	TX	62390
5	NY	51546
6	MN	49256
7	VA	44853
8	NC	43935
9	PA	40541
10	IL	32635
11	MD	29750
12	AZ	26139
13	SC	25944
14	UT	24882
15	NJ	24612
16	GA	23822
17	MI	20846
18	TN	19224
19	LA	18204
20	OH	17948

**Top 20 states that has the highest number of accidents in the US**



I think this dataset is representative of the general distribution of accidents in the US since it has a lot of data points and it covers nearly every state in the US. This dataset contains information about 49 states and there are 50 states in the US. One thing I notice is that California has much more accidents than all the other states. In general, I think this dataset reflects the distribution of accidents in the US.